



Problems with Language Modelling

Rene Pickhardt, Heinrich Hartmann



Outline

- 1) Basics / Background
- 2) Conjectures
- 3) Justification for the strong assumption
- 4) Preliminary experimental results
- 5) Smoothing methods
- 6) Summary

Basics & Background

A Language Model assigns each sentence a probability

$$P_{LM} : S \longrightarrow [0, 1] \qquad S := W^*$$

$$P_{LM}(s) = p_s \qquad W := \{w_1, w_2, \dots, w_N\}$$

- Probability means the following equation must hold

$$\sum_{s \in S} P_{LM}(s) = 1$$

- Example:

$$P_{LM}(\textit{JOHN READ MOBY DICK}) = 7.39 * 10^{-15}$$

Not to confuse with an n-gram model

$$P^n : W^n \longrightarrow [0, 1]$$

$$W^n := \underbrace{W \times \dots \times W}_{n\text{-times}}$$

$$P^n(s) = p_s$$

$$W := \{w_1, w_2, \dots, w_N\}$$

- Probability means the following equation must hold

$$\sum_{w_1^n \in W^n} P^n(w_1^n) = 1 \quad w_1^n := w_1 \dots w_n = (w_1, \dots, w_n) \in W^n$$

- Example:

$$P^4(\text{JOHN READ MOBY DICK}) = 9.23 * 10^{-9}$$

$$P^5(\text{JOHN READ MOBY DICK}) = \emptyset$$

When doing Computing we usually do not have LMs

- Even when you see $P(s)$ people can only calculate $P^n(s)$
 - This has some strong impact on the semantics of the formulas in the field (\rightarrow later in this talk)
- Even over a finite set of words the set of all sentences is of infinite size
- N-gram models only estimate Language models
 - ♦ Together with Heinrich we can:
<https://github.com/HeinrichHartmann/doc/blob/a6f168119d36ac5e6766bef2e34322b907e8bc4f/LanguageModels.pdf>
 - Construct n-gram model from Language models
 - Construct Language models from n-gram models
 - Switch between models of various length
 - All just theoretically without doing any statistics and counting

Impact of language models for applications

- Noisy channel model

$$w' = \underset{w \in W}{\operatorname{argmax}} \underbrace{P_{\text{domain}}(O|w)}_{\text{likelihood}} \overbrace{P_{LM}(w)}^{\text{prior}}$$

- The prior is the Language Model.
- More precise language models yield better applications
- Obviously $P_{LM}(w)' = cP_{LM}(w)$ with $c \neq 1$ will not change the application (w' will be the same)
- Take care with the formula it is not clear if w is just a word!

Language Models are evaluated using entropy

- Entropy is defined as

$$H(P_{LM}) = - \sum_{t \in T} P_{MLE}(t) \log(P_{LM}(t))$$

- T is a Test corpus with test sequences t
- The better the Language model predicts the Testcorpus the smaller the entropy \rightarrow small entropy values are great

$$P'(t) = cP(t), c > 1 \longrightarrow \sum_{s \in S} P(s) = c \neq 1$$

- But

$$H(P') < H(P)$$

Chen 1998: Entropy is not a good measure for Language Modeling

- Use LMs (n-gram Models) for applications like
 - speech recognition (Used in the original paper)
 - machine translation
 - spell checking
 - autocompletion
- Entropy of complexly build language models does not correlate with Metrics in applications
- This is somehow community knowledge

Common knowledge in the Community

- Quote Chen 1998 (Evaluation metrics for Language Models):
“While perplexities can be calculated efficiently and without access to a speech recognizer, they often do not correlate well with speech recognition word-error rates.”
- Alfonseca (Google) in reply to our proposal:
“As you know, it is not always the case that a decrease in perplexity leads to improvements in an extrinsic evaluation, be it machine translation, speech, or something else”
- Stupid backoff (Brants et. al. 2007)

Conjecture

Entropy is an excellent measure for Language Modeling

(and it correlates with metrics from applications)

- Reformulation: The Language Models used in the Chen paper and also later have not been proper probability functions. When going to probability functions entropy will work best.
- We will see later how we can justify this bold claim:
 - In implementation most LMs (n-gram models) are currently not a proper probability function


More questions:

- Can we fix the existing Language Models?
- What happens to the applications after fixing the LMs?
- Might applications with unfixed models still outperform fixed models?

More questions:

- Can we fix the existing Language Models?
 - (We strongly believe: yes in all cases)
- What happens to the applications after fixing the LMs?
 - (We expect: that the applications perform better)
- Might applications with unfixed models still outperform fixed models?
 - If no: perfect! Claim made
 - If yes: We should stop putting effort in building probability functions. (Google already stopped)

From Brants et. al. 2007: Stupid back off



$$S(w_i|w_{i-k+1}^{i-1}) =$$

$$\begin{cases} \frac{f(w_{i-k+1}^i)}{f(w_{i-k+1}^{i-1})} & \text{if } f(w_{i-k+1}^i) > 0 \\ \alpha S(w_i|w_{i-k+2}^{i-1}) & \text{otherwise} \end{cases}$$

In general, the backoff factor α may be made to depend on k . Here, a single value is used and heuristically set to $\alpha = 0.4$ in all our experiments.

The value of 0.4 was chosen empirically based on good results in earlier experiments. Using multiple values depending on the n -gram order slightly improves results.

Figure 5: BLEU scores for varying amounts of data using Kneser-Ney (KN) and Stupid Backoff (SB).

My interpretation of Brants

- Stupid backoff is certainly not a probability function.
 - ◆ (Not even in theory)
- For large scale experiments the results are getting closer since almost no backoff steps take place
- Models work better on large scale data not only because of sparsity but also because the implementation mistakes that are being made are becoming less and less
 - ◆ (large data \Rightarrow no backoff necessary)

Justification of Assumptions

Let us assume a Maximum likelihood n-gram model

$$P_{MLE}^n(w_1^n) := \frac{c(w_1^n)}{c(-\dots-)}$$

with

$$\sum_{w_1^n \in W^n} c(w_1^n) =: c(\underbrace{-\dots-}_{n\text{-times}})$$

We test for probability:

$$\sum_{w_1^n \in W^n} P_{MLE}^n(w_1^n) = \frac{\sum_{w_1^n \in W^n} c(w_1^n)}{c(-\dots-)} = \frac{c(-\dots-)}{c(-\dots-)} = 1$$

We have the following conditional (?) probabilities

$$P_{MLE-sem}^n(w_n | w_1^{n-1}) = \frac{c(w_1^n)}{\sum_{w_n \in W} c(w_1^{n-1} w_n)} = \frac{c(w_1^n)}{c(w_1^{n-1} -)}$$

- We often see the following implementation

$$P_{MLE-impl}^n(w_n | w_1^{n-1}) = \frac{c(w_1^n)}{c(w_1^{n-1})}$$

- But (!)obviously

$$\frac{c(w_1^n)}{c(w_1^{n-1})} \neq \frac{c(w_1^n)}{c(w_1^{n-1} -)}$$

- The difference over all words can just be measured as the number of sentences in the corpus

Chain rule (probability)

From Wikipedia, the free encyclopedia

In **probability** theory, the **chain rule** permits the calculation of any member of the **joint distribution** of a set of **random variables** using only **conditional probabilities**. The rule is useful in the study of **Bayesian networks**, which describe a probability distribution in terms of conditional probabilities.

Consider an indexed set of sets A_1, \dots, A_n . To find the value of this member of the joint distribution, we can apply the definition of conditional probability to obtain:

$$P(A_n, \dots, A_1) = P(A_n | A_{n-1}, \dots, A_1) \cdot P(A_{n-1}, \dots, A_1)$$

Repeating this process with each final term creates the product:

$$P\left(\bigcap_{k=1}^n A_k\right) = \prod_{k=1}^n P\left(A_k \mid \bigcap_{j=1}^{k-1} A_j\right)$$

With four variables, the chain rule produces this product of conditional probabilities:

$$P(A_4, A_3, A_2, A_1) = P(A_4 | A_3, A_2, A_1) \cdot P(A_3 | A_2, A_1) \cdot P(A_2 | A_1) \cdot P(A_1)$$

References: https://www.ibm.com/developerworks/community/blogs/nlp/entry/the_chain_rule_of_probability

Let us have a look at the chain rule of probability

$$P(w_1^n) := P(w_n | w_1^{n-1}) P(w_{n-1} | w_1^{n-2}) \dots P(w_2 | w_1) P(w_1)$$

- No Problem for a real language Model but for n-gram models we run into Problems

$$P^n(w_1^n) := P^n(w_n | w_1^{n-1}) P^?(w_{n-1} | w_1^{n-2}) \dots P^?(w_1)$$

- There are various interpretations for the “?”
 - First choice decrease n \rightarrow (D = decrease)

$$P_D^n(w_1^n) := P_D^n(w_n | w_1^{n-1}) P_D^{n-1}(w_{n-1} | w_1^{n-2}) \dots P_D^1(w_1)$$

- But there is a problem (no prove here but some data later)

$$P_D^n(w_1^n) \neq P_{MLE}^n(w_1^n)$$

Introduce skips to rescue the chain rule

$$P^n(w_1^n) := P^n(w_n | w_1^{n-1}) P^?(w_{n-1} | w_1^{n-2}) \dots P^?(w_1)$$

Define

$$P^n(w_{n-l} | w_1^{n-1-l}) := \frac{c(w_1^{n-l} \overbrace{\dots}^{l\text{-times}})}{c(w_1^{n-1-l} \underbrace{\dots}_{(l+1)\text{-times}})}$$

With this definition we can easily see (prove blackboard) that:

$$P^n(w_1^n) = P_{MLE}^n(w_1^n)$$

Preliminary test results

Our test corpora

- Test 1 (3 Sent, 3 tokens, 18 words)
 - ♦ a b c a a b
 - ♦ b a a b c a
 - ♦ c a b a b a
- Moby Dick Corpus (3 Sent, 11 tokens, 15 words)
 - JOHN READ MOBY DICK
 - MARY READ A DIFFERENT BOOK
 - SHE READ A BOOK BY CHER
- Wikipedia (very small only entropy)
 - (6k Sent, ??? tokens, ~50 k words)

Only MLE with skips is a proper probability distribution

Sum of probabilities for all possible n-grams of various estimated n-gram models

n	SIKP	D	impl
1	1.00	1.00	1.00
2	1.00	1.00	0.83
3	1.00	1.00	0.67
4	1.00	1.00	0.50
5	1.00	0.82	0.33

ABC Corpus

	SKIP	D	impl
1	1.00	1.00	1.00
2	1.00	0.87	0.80
3	1.00	0.67	0.60
4	1.00	0.40	0.40
5	1.00	0.20	0.20

Moby Dick Corpus

	SKP	SRILM
1	1	1
2	1	0.0037

Wiki corpus (6k sentences)

P_D is not a probability function

- Problem (rather subtle):
 - We divide 0/0 and assign this probability 0
 - Double check if current libs make this error too (We guess yes because it is very easy to make that mistake)

What does the following distribution look like?


$$P(w|CHER) = \frac{c(CHER, w)}{c(CHER_)} = \frac{0}{0} = ?$$

JOHN READ MOBY DICK
MARY READ A DIFFERENT BOOK
SHE READ A BOOK BY CHER

ABC Corpus sum of probabilities vs Entropy gain


ABC Corpus sum of probabilities vs Entropy gain

remember entropy gains are bad



Moby Dick Corpus Entropy in comparison to mistake


Sum over all probabilities vs relative entropy



Difference in Entropy (attention |T| depends on n) (wiki)

Relative Change in Entropy

difference of a correct model to two common implementations



$$y = H(P_{MODEL}) - H(P_{MLE})$$

ABC: Entropy & sum of 'probabilities'; different estimator

	SUM	H
1	1.00	0.52
2	1.00	0.73
3	1.00	0.93
4	1.00	0.92
5	1.00	0.78

P_{MLE}

	SUM	H
1	1.00	0.52
2	1.00	0.73
3	1.00	0.92
4	1.00	0.92
5	0.82	0.88

P_D

	SUM	H
1	1.00	0.52
2	0.83	0.80
3	0.67	1.10
4	0.50	1.22
5	0.33	1.26

P_{impl}

Note: the test corpus consisted of all seen sequences
(in any other case entropy would be infinity)

Mini conclusion

- Entropy values are wrong to the disadvantage of the community
 - (on large corpora the mistake seems negligible)
- Implementation P_{impl} and fixed implementation P_D don't produce a probability functions
 - This can be fixed for P_D
- Introducing skips does at least make the Math correctly
 - and leads right a way to the notion of generalized language models
- We have to make more experiments with standard libs

Smoothing methods

Existing Smoothing methods

- Unigram:
 - ◆ Good Turing
 - ◆ Laplace
- Advanced:
 - ◆ Absolute discounting
 - ◆ Backoff methods
 - ◆ Interpolation methods
 - ◆ Kneser Ney Smoothing (Absolute discounting)
 - ◆ Modified Kneser Ney Smoothing
 - ◆ Witten Bell Smoothing
 - ◆ Katz Smoothing

Ideas and goals of smoothing

- Idea:
 - ◆ Given probability distribution
 - ◆ Make it more 'even'
 - ◆ Lower large probabilities
 - ◆ Increase low / zero probabilities
- Goal for language modeling
 - ◆ No probability should be 0
 - (that might be a hard and also stupid goal)
- Goal is needed for applications not to break

Problems with smoothing

- In theory those smoothing methods work greatly
 - ◆ (proves can be received from Heinrich)
- In implementations there are many sources for adding too much probability (similar to P_D)
 - ◆ Sum of all probabilities > 1
 - ◆ Probably leading to an Entropy drop
 - (which should not be there)
- Chen, 1998: “Interpol always works better than back off”
 - ◆ WE: Backoff is certainly a probability function
 - ◆ Interpol implementations make backoffs when they are not supposed to (again 0/0 is the problem)

Summary

There is still a lot to do

- Check all implementations of all toolkits
- Disprove Chen (Entropy is a good measure)
- Check how big the effects really become in applications and on real size data sets
- Fix implementations (give hints what the semantics are and what can be done)
- Provide a full mathematical framework / theory of Language modeling and smoothing methods

The community has some “magic tricks”

- Including BOS and EOS tokens to sentences
- Include UNK tag for unknown words
- Smoothing unigram distributions
- → We have to see which of these address the above mentioned problems to which extend